

# Phase 1 · EXP-A1 — Baseline: No Ambient (Track A, ODAS-free)

**experiments.pdf** tag: `exp_a1_no_ambient`. **Question (Phase 1):** does ambient type drive ODAS false positives, and does matched-ambient training help? EXP-A1 is the **noise-floor baseline** — directional sources only, no ambient.

**Scope of this run — Track A (ODAS-free):** produced `model_a1_gt` (the GT-trained arm) entirely on this Mac: scene → render → GT clips → frozen-YAMNet head → clip accuracy. The **post-ODAS arm** (`model_a1_odas`) and the FP/min metric require the `odaslive` binary (Track B) and are deferred to a Linux/Pi env — see [Status & deferred](#).


## What ran

| Step                                                            | Tool                                                                       | Output                                                                       |
|-----------------------------------------------------------------|----------------------------------------------------------------------------|------------------------------------------------------------------------------|
| Scene (6 labels, 600 s, dist 10–150 m, <b>0 ambient</b> )       | <code>expA1_render.py</code> (headless <code>renderer.py</code> )          | 90 directional sources, <code>.raw</code> + 90 <code>.f32</code> GT sidecars |
| GT clip dataset (3 s win, 1.5 s hop, source-level split)        | <code>expA1_build_dataset.py</code> ( <code>gt_dataset_builder.py</code> ) | <code>gt_a1_no_ambient</code> : 133 clips, 6 labels                          |
| Holdout (300 s, seed 99)                                        | same, <code>--name exp_a1_holdout</code>                                   | <code>gt_a1_holdout</code> : 111 clips                                       |
| <code>model_a1_gt</code> (frozen YAMNet + 256→128→softmax head) | <code>expA1_train_gt.py</code>                                             | clip-acc + confusion                                                         |

Fixed protocol honored: room 250×250×20, absorption 0.7, max\_order 3, ReSpeaker 4-mic geometry, 16 kHz. Labels: Elephant, Bear, Frog, Lion, drone\_bebop, drone\_binary.

## Result

| Metric                              | Value                                                                               |
|-------------------------------------|-------------------------------------------------------------------------------------|
| Internal test clip accuracy         | <b>0.71–0.81</b> (21-clip fold, run-to-run variance)                                |
| <b>Holdout-render clip accuracy</b> | <b>0.60</b>                                                                         |
| Holdout per-class F1                | Elephant 0.84, Frog 0.84, drone_binary 0.48, Lion 0.45, Bear 0.38, drone_bebop 0.35 |



## Findings

- **A clean baseline exists and the chain works end-to-end** on this Mac (render → GT dataset → trained model → eval), with real artifacts under [experiments/sim/](#).
- **~20 pt generalization drop** from internal test (0.81) to a separate holdout render (0.60), *even with no ambient and clean GT audio*. The gap is driven by **Bear, drone\_bebop, Lion, drone\_binary; Elephant and Frog** hold up ( $\approx 0.84$ ). This matches the PDF's observation that short/quiet events (Bear) are disproportionately hard, and tells us the floor is set by intra-class acoustic variability + tiny per-class counts, before ODAS even enters.
- **Drone sub-types confuse each other** (drone\_binary ↔ drone\_bebop) — expected, since both are motor/rotor spectra; per-class thresholds (Phase 3 C2) will matter here.

## Caveats

- **Small dataset** (133 train / 111 holdout clips). The instance counts were modest to keep renders fast; the headline number is the *holdout* (comparable across experiments), and the internal-test figure is noisy. Scale [--instances](#) for tighter numbers.
- **Track A ≠ full deployment metric**. EXP-A1's headline metric in the PDF is **FP/min ≤ 2**, which is intrinsically an ODAS-detection metric (detections during quiet periods). It cannot be computed without [odaslive](#). Track A gives clip accuracy only.

- **Head, not full 2-phase finetune.** This run uses the frozen-backbone Phase-1 head. The upstream `train_yamnet.py` Phase-2 (top-4 backbone unfreeze) currently breaks rebuilding the backbone from the `tensorflow/models` layer-defs (Keras-version / `_YAMNET_LAYER_DEFS` calling-convention mismatch vs current `master`). The GT-vs-ODAS comparison the experiments test is **unaffected** — the variable is the training *data*, and both arms use the identical head procedure — but the exact deployable TFLite needs the pinned `tensorflow/models` commit in the Linux env.

Status & deferred

| EXP-A1 deliverable                                     | Status                                  |
|--------------------------------------------------------|-----------------------------------------|
| Scene + render + GT dataset                            | ✔ done (this Mac)                       |
| <code>model_a1_gt</code> + clip accuracy               | ✔ done                                  |
| <code>model_a1_odas</code> (post-ODAS curator dataset) | ⏸ needs <code>odaslive</code> (Track B) |
| Event P/R + <b>FP/min</b> on holdout                   | ⏸ needs <code>odaslive</code> (Track B) |

How to reproduce

```
.venv/bin/python experiments/scripts/expA1_render.py --duration 600 --instances 15 --seed 1 --name exp_a1_no_ambient
.venv/bin/python experiments/scripts/expA1_build_dataset.py --meta experiments/sim/renderers/renderers/exp_a1_no_ambient_*.json --dataset gt_a1_no_ambient
.venv/bin/python experiments/scripts/expA1_render.py --duration 300 --instances 10 --seed 99 --name exp_a1_holdout
.venv/bin/python experiments/scripts/expA1_build_dataset.py --meta experiments/sim/renderers/renderers/exp_a1_holdout_*.json --dataset gt_a1_holdout
.venv/bin/python experiments/scripts/expA1_train_gt.py --dataset experiments/sim/gt_datasets/gt_datasets/gt_a1_no_ambient --holdout experiments/sim/gt_datasets/gt_datasets/gt_a1_holdout
```

Next

EXP-A2/A3 (synthetic Rain/Wind ambient) reuse this exact harness — just add ambient sources to the scene generator. The post-ODAS arm and FP/min for all of Phase 1 unlock once `odaslive` is built (Linux/Pi). See `08_deployment_and_loop.md` for the ODAS build.